Group 6 - Cancer Predictions

Raima Ghosh - Shane Abbley - Janell Napper - Amit Choksi




Which lifestyle choices determine the severity of a cancer diagnosis?

Data and Data Cleaning

- We obtained our lung cancer data from Kaggle.
- Our data contains symptom severity and lifestyle choices ranked from 1 to 8 and cancer severity ranked as either high, moderate, or low.
- We cleaned the data using pandas in a jupyter notebook.
  - We checked for empty data.
  - We dropped all of the symptom columns and converted them to a separate data frame for our visualizations.
  - We converted the cancer severity to a numeric value.

Patient Id	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease	Balanced Diet	Obesity	Smoking	Passive Smoker	Cancer Severity
P1	33	1	2	4	5	4	3	2	2	4	3	2	1
P10	17	1	3	1	5	3	4	2	2	2	2	4	2
P100	35	1	4	5	6	5	5	4	6	7	2	3	3
P1000	37	1	7	7	7	7	6	7	7	7	7	7	3
P101	46	1	6	8	7	7	7	6	7	7	8	7	3
P102	35	1	4	5	6	5	5	4	6	7	2	3	3
P103	52	2	2	4	5	4	3	2	2	4	3	2	1
P104	28	2	3	1	4	3	2	3	4	3	1	4	1

Database Construction and Entity Relationship Diagram

- We wanted a locally hosted database and our dataset is relatively small, so we chose to use a SQLite database.
- A simple way to create a SQLite database is to use the sqlite3 module in python.
- After creating a database, queries can be run with SQL syntax.
- A table was created for both the symptoms and the lifestyle choices.
- A combined table was created using a full outer join of the lifestyle and symptoms tables.

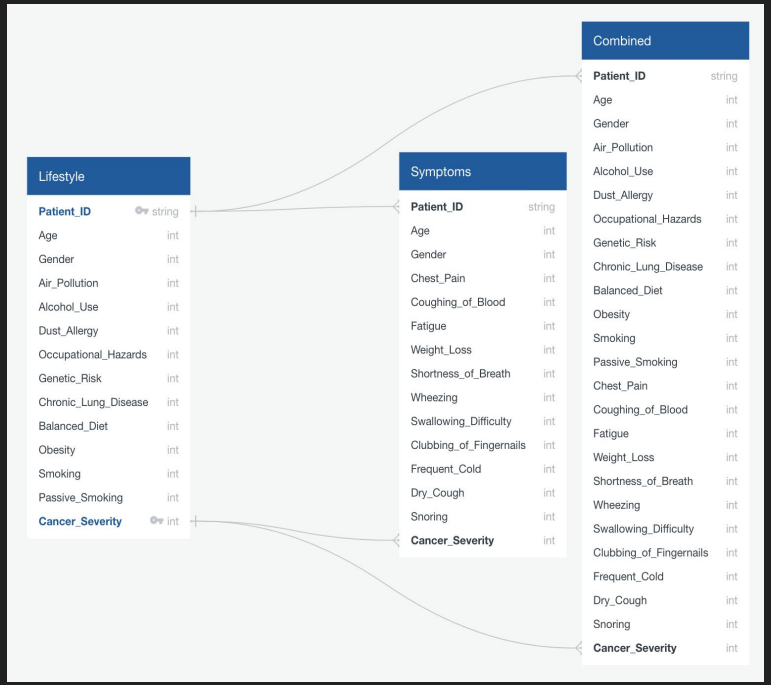

Entity relationship diagram for our database

After looking over this data, we wanted to answer these three questions.

1. Which lifestyle choices are most associated with a higher severity of cancer?
2. How can we most accurately predict cancer severity using machine learning?
3. Which machine learning model predicts the severity of cancer most accurately?

Cancer Severity by Level

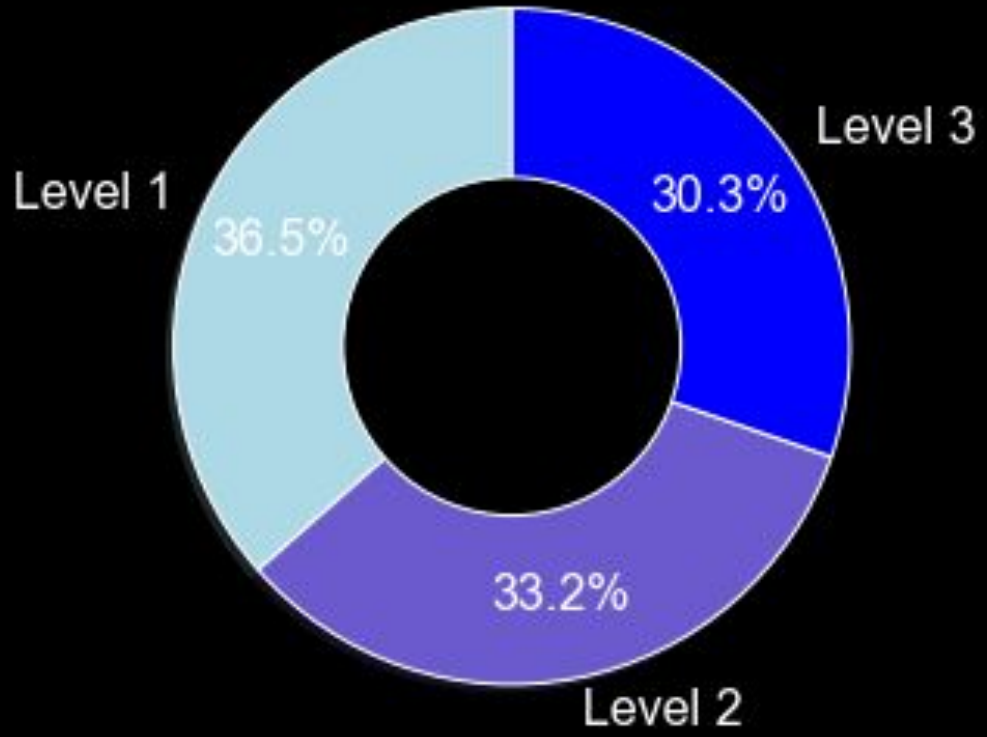

Top Three Columns with the Highest usage value (8)

```
In [65]: cancer_patient_df["Alcohol use"].value_counts()

Out[65]: 2    202
         8    188
         7    167
         1    152
         5     90
         3     80
         6     80
         4     41
         Name: Alcohol use, dtype: int64
```

```
In [75]: cancer_patient_df["Passive Smoker"].value_counts()

Out[75]: 2    284
         7    187
         4    161
         3    140
         8    108
         1     60
         6     30
         5     30
         Name: Passive Smoker, dtype: int64
```

```
In [74]: cancer_patient_df["Smoking"].value_counts()

Out[74]: 2    222
         7    207
         1    181
         3    172
         8     89
         6     60
         4     59
         5     10
         Name: Smoking, dtype: int64
```

Cancer Severity Per Lifestyle Choice

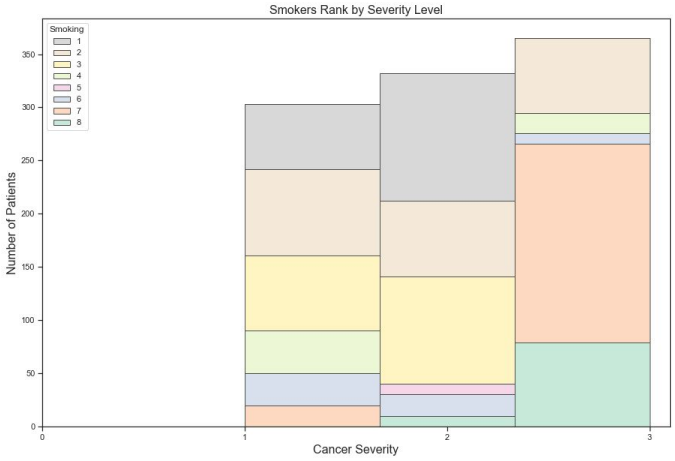

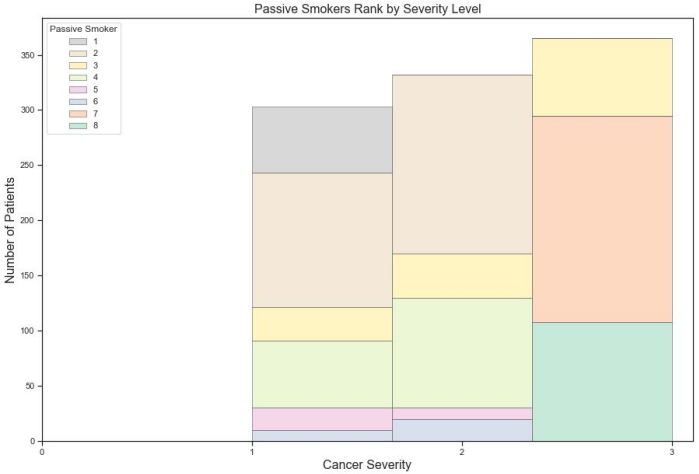

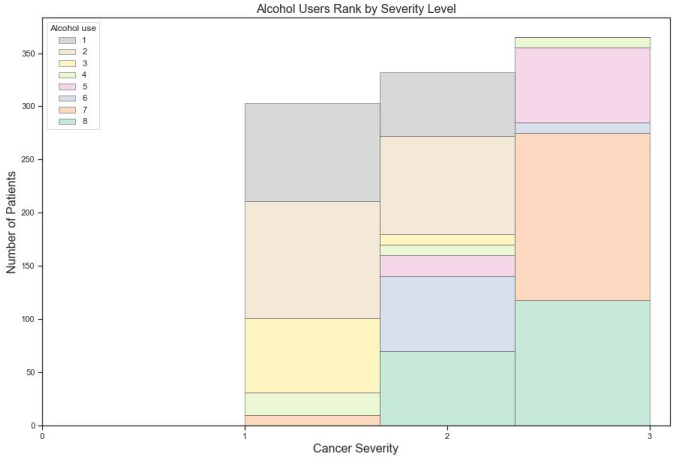

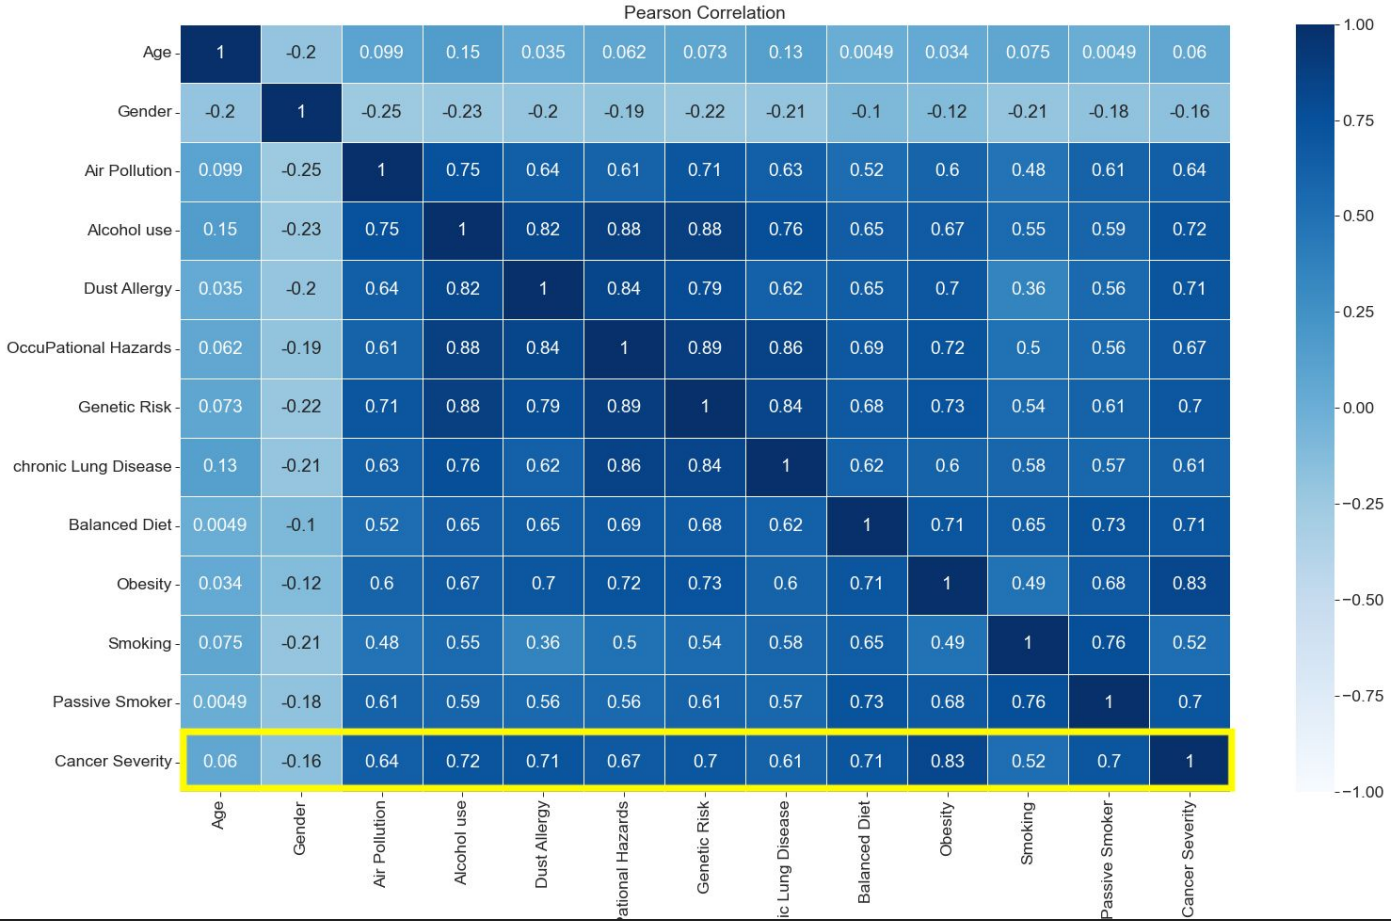
Pearson Correlation
Age 1 -0.2 0.099 0.15 0.035 0.062 0.073 0.13 0.0049 0.034 0.075 0.0049 0.06
Gender -0.2 1 -0.25 -0.23 -0.2 -0.19 -0.22 -0.21 -0.1 -0.12 -0.21 -0.18 -0.16
Air Pollution 0.099 -0.25 1 0.75 0.64 0.61 0.71 0.63 0.52 0.6 0.48 0.61 0.64
Alcohol use 0.15 -0.23 0.75 1 0.82 0.88 0.88 0.76 0.65 0.67 0.55 0.59 0.72
Dust Allergy 0.035 -0.2 0.64 0.82 1 0.84 0.79 0.62 0.65 0.7 0.36 0.56 0.71
OccuPational Hazards 0.062 -0.19 0.61 0.88 0.84 1 0.89 0.86 0.69 0.72 0.5 0.56 0.67
Genetic Risk 0.073 -0.22 0.71 0.88 0.79 0.89 1 0.84 0.68 0.73 0.54 0.61 0.7
chronic Lung Disease 0.13 -0.21 0.63 0.76 0.62 0.86 0.84 1 0.62 0.6 0.58 0.57 0.61
Balanced Diet 0.0049 -0.1 0.52 0.65 0.65 0.69 0.68 0.62 1 0.71 0.65 0.73 0.71
Obesity 0.034 -0.12 0.6 0.67 0.7 0.72 0.73 0.6 0.71 1 0.49 0.68 0.83
Smoking 0.075 -0.21 0.48 0.55 0.36 0.5 0.54 0.58 0.65 0.49 1 0.76 0.52
Passive Smoker 0.0049 -0.18 0.61 0.59 0.56 0.56 0.61 0.57 0.73 0.68 0.76 1 0.7
Cancer Severity 0.06 -0.16 0.64 0.72 0.71 0.67 0.7 0.61 0.71 0.83 0.52 0.7 1
Age
Gender
Air Pollution
Alcohol use
Dust Allergy
Genetic Risk
Balanced Diet
Obesity
Smoking
Passive Smoker
Cancer Severity
1.00
0.75
0.50
0.25
0.00
−0.25
−0.50
−0.75
−1.00

Model Overview

- Model of Choice : Logistic Regression & SVM for the predicting the severity of cancer as low, medium, or high
- We can compare the various results from running these algorithms
- Ran the model with different solvers, and different values of iterations.
- Logistic regression : This algorithm is used for classification problems in machine learning.
- Support vector machine : This algorithm separates the data points using a line , this line is chosen such that it will be furthermost from the nearest data points in 2 categories.

Logistic Regression

It is a classification model which is used to predict the odds in favour of a particular event. The odds ratio represents the positive event which we want to predict, for example, how likely a sample has cancer/ how likely is it for an individual to become diabetic in future. It used the sigmoid function to convert an input value between 0 and 1. It can further be extended to multiple logistic regression.

Logistic Regression tries to maximize the conditional likelihood of the training data, it is highly prone to outliers. Standardization (as co-linearity checks) is also fundamental to make sure a features' weights do not dominate over the others.

Source :
https://www.geeksforgeeks.org/differentiate-between-support-vector-machine-and-logistic-regression/

SVM

A powerful classification algorithm for predicting classification problems. To maximize the margin among class variables. This margin (support vector) represents the distance between the separating hyperplanes (decision boundary). The reason to have decision boundaries with large margin is to separate positive and negative hyperplanes with adjustable bias-variance proportion.

Source : https://www.geeksforgeeks.org/differentiate-between-support-vector-machine-and-logistic-regression/

Comparison Pros/Cons of using various models

Benefits of Logistic regression

1. Solving classification problem
2. Works with already identified independent variable

Cons

1. Vulnerable to overfitting
2. Can miss outliers / lower sensitivity for large unbalanced data sets

SVM

1. Tries to find the best margin that separates the classes that reduces the risk of error on the data
2. Risk of overfitting is less.
3. Gave us the best prediction for the model under study.

Logistic Regression (liblinear)

Logistic Regression Solver(liblinear)

	Predicted Low	Predicted Medium	Predicted High
Actual Low	69	26	0
Actual Medium	21	81	8
Actual High	0	0	95

Accuracy Score: 0.8166666666666667

Classification Report

	precision	recall	f1-score	support
1	0.77	0.73	0.75	95
2	0.76	0.74	0.75	110
3	0.92	1.00	0.96	95
accuracy			0.82	300
macro avg	0.82	0.82	0.82	300
weighted avg	0.81	0.82	0.81	300

Logistic Regression libfgs

Logistic Regression Solver(libfgs)

	Predicted Low	Predicted Medium	Predicted High
Actual Low	77	18	0
Actual Medium	18	88	4
Actual High	0	0	95

Accuracy Score: 0.8666666666666667

Classification Report

	precision	recall	f1-score	support
1	0.81	0.81	0.81	95
2	0.83	0.80	0.81	110
3	0.96	1.00	0.98	95
accuracy			0.87	300
macro avg	0.87	0.87	0.87	300
weighted avg	0.86	0.87	0.87	300

Logistic Regression newton-cg

Logistic Regression Solver(newton-cg)

	Predicted Low	Predicted Medium	Predicted High
Actual Low	77	18	0
Actual Medium	18	88	4
Actual High	0	0	95

Accuracy Score: 0.8666666666666667

Classification Report

	precision	recall	f1-score	support
1	0.81	0.81	0.81	95
2	0.83	0.80	0.81	110
3	0.96	1.00	0.98	95
accuracy			0.87	300
macro avg	0.87	0.87	0.87	300
weighted avg	0.86	0.87	0.87	300

Logistic Regression sag

Logistic Regression Solver(sag)

	Predicted Low	Predicted Medium	Predicted High
Actual Low	72	23	0
Actual Medium	18	88	4
Actual High	0	0	95

Accuracy Score: 0.85

Classification Report

	precision	recall	f1-score	support
1	0.80	0.76	0.78	95
2	0.79	0.80	0.80	110
3	0.96	1.00	0.98	95
accuracy			0.85	300
macro avg	0.85	0.85	0.85	300
weighted avg	0.85	0.85	0.85	300

Logistic Regression saga

Logistic Regression Solver(saga)

	Predicted Low	Predicted Medium	Predicted High
Actual Low	72	23	0
Actual Medium	18	84	8
Actual High	0	0	95

Accuracy Score: 0.8366666666666667

Classification Report

	precision	recall	f1-score	support
1	0.80	0.76	0.78	95
2	0.79	0.76	0.77	110
3	0.92	1.00	0.96	95
accuracy			0.84	300
macro avg	0.84	0.84	0.84	300
weighted avg	0.83	0.84	0.83	300

SVG Algorithm (Best prediction)

SVG

	Predicted Low	Predicted Medium	Predicted High
Actual Low	76	19	0
Actual Medium	18	92	0
Actual High	0	0	95

Accuracy Score: 0.8766666666666667

Classification Report

	precision	recall	f1-score	support
1	0.81	0.80	0.80	95
2	0.83	0.84	0.83	110
3	1.00	1.00	1.00	95
accuracy			0.88	300
macro avg	0.88	0.88	0.88	300
weighted avg	0.88	0.88	0.88	300

Dashboard

Early View of the Dashboard

- Coded using d3.json with Bootstrap components
- Csv files will be cleaned using pandas and converted to json
- Will also display features that come out of ML model and preliminary data graphs
- Interactive input connected to Symptom info and bar graph

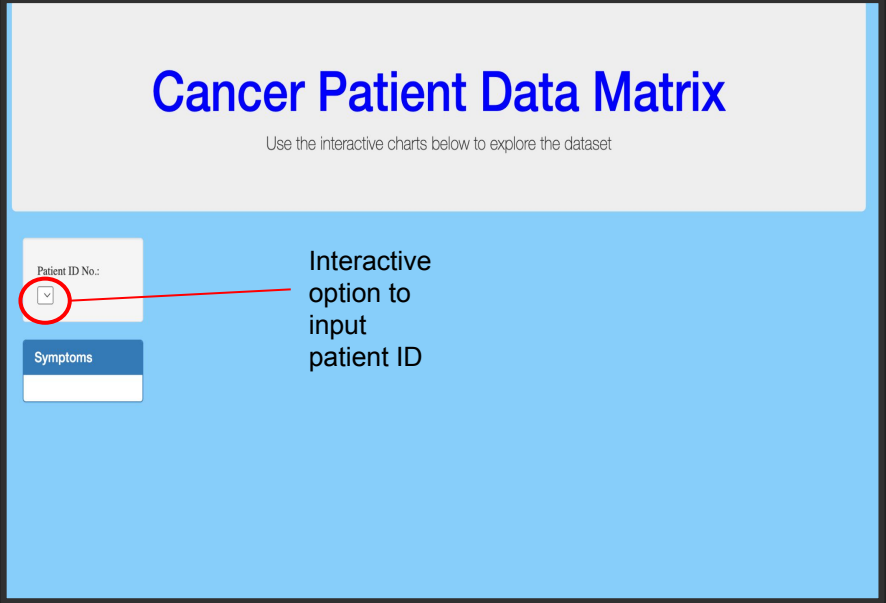

Interactive option to input patient ID

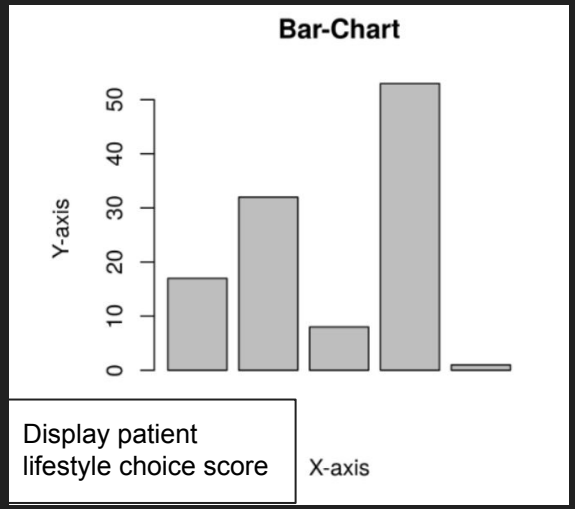

Display patient lifestyle choice score

Blueprint

Classify lifestyle scores at different cancer severity

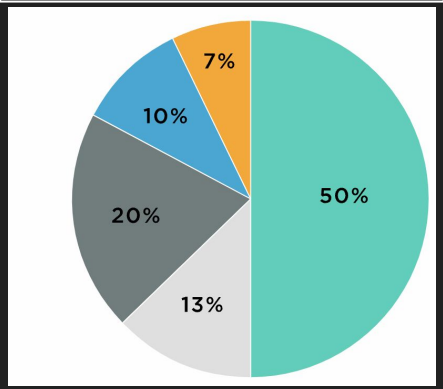

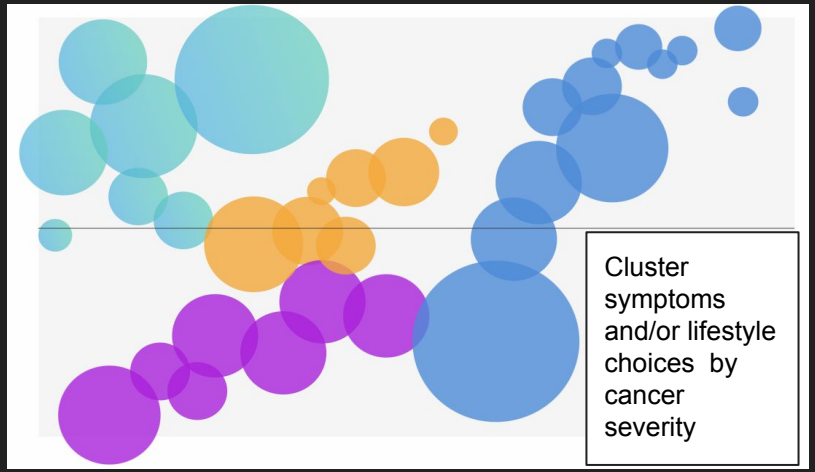

Cluster symptoms and/or lifestyle choices by cancer severity

Dashboard

Final View

Interactive option to input patient ID

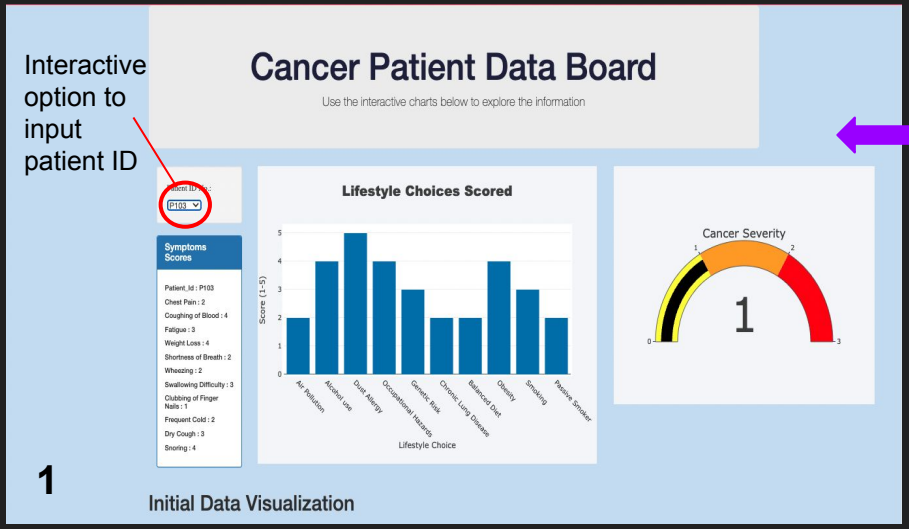

1

Initial Data Visualization

When selected, displays:

- Symptoms scores
- Lifestyle scores bar graph
- Cancer severity gauge

Classifies 3 hypothesized lifestyle habits by different cancer severity

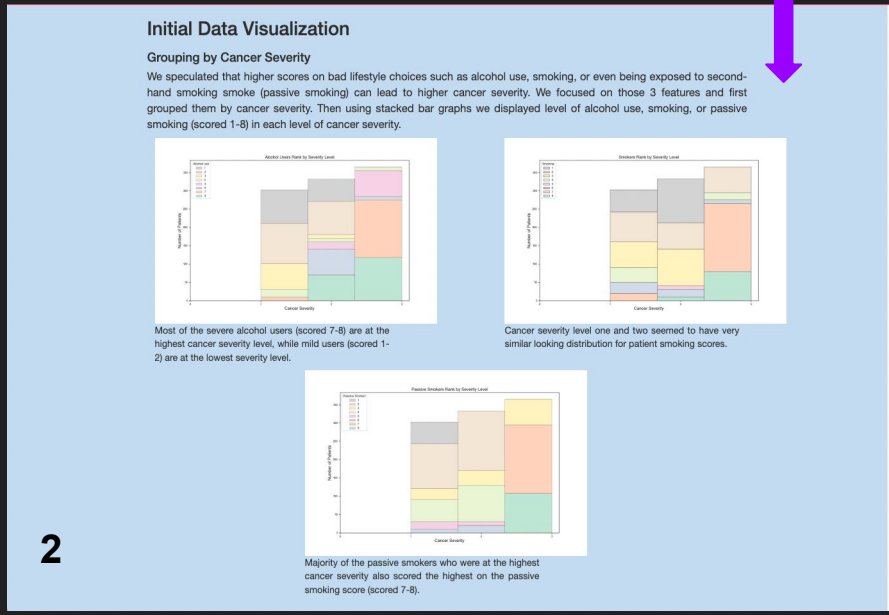

Current View of the Dashboard

- Coded using **d3.json** with **CSS** styling and **Bootstrap** components
- Dashboard uses **HTML**

Dashboard

Final View

Bubble chart displays the **feature importance**

3

Outcomes from Machine Learning Model

Feature Importance Outcomes from the Machine Learning Models

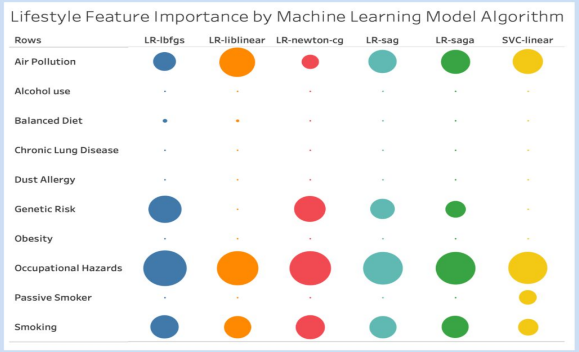

We ran multiple machine learning models on our dataset. As summarized in the bubble chart (above), we see that "Occupational Hazard" is a highly ranked feature, followed by "Air Pollution", "Smoking", and "Genetic Risk". Out of all the features we hypothesized previously would be top ranked, the machine learning models only highlighted "Smoking".

Confusion matrices from the 6 ML models displayed as table

4

Confusion Matrices Generated from the Machine Learning Models

LR-liblinear	Predicted Low	Predicted Medium	Predicted High
Actual Low	69	26	0
Actual Medium	21	81	8
Actual High	0	0	95

LR-lbfgs	Predicted Low	Predicted Medium	Predicted High
Actual Low	77	18	0
Actual Medium	18	88	4
Actual High	0	0	95

Newton-cg	Predicted Low	Predicted Medium	Predicted High
Actual Low	77	18	0
Actual Medium	18	88	4
Actual High	0	0	95

LR-sag	Predicted Low	Predicted Medium	Predicted High
Actual Low	72	23	0
Actual Medium	18	88	4
Actual High	0	0	95

LR-saga	Predicted Low	Predicted Medium	Predicted High
Actual Low	72	23	0
Actual Medium	18	84	8
Actual High	0	0	95

SVC-linear	Predicted Low	Predicted Medium	Predicted High
Actual Low	76	19	0
Actual Medium	18	92	0
Actual High	0	0	95

The table above displays the confusion matrix for all the machine learning models that were run over the course of our analysis. From the outcomes above we see that SVC model has the best results to predict the highest cancer severity.

Dashboard

Final View- Demo

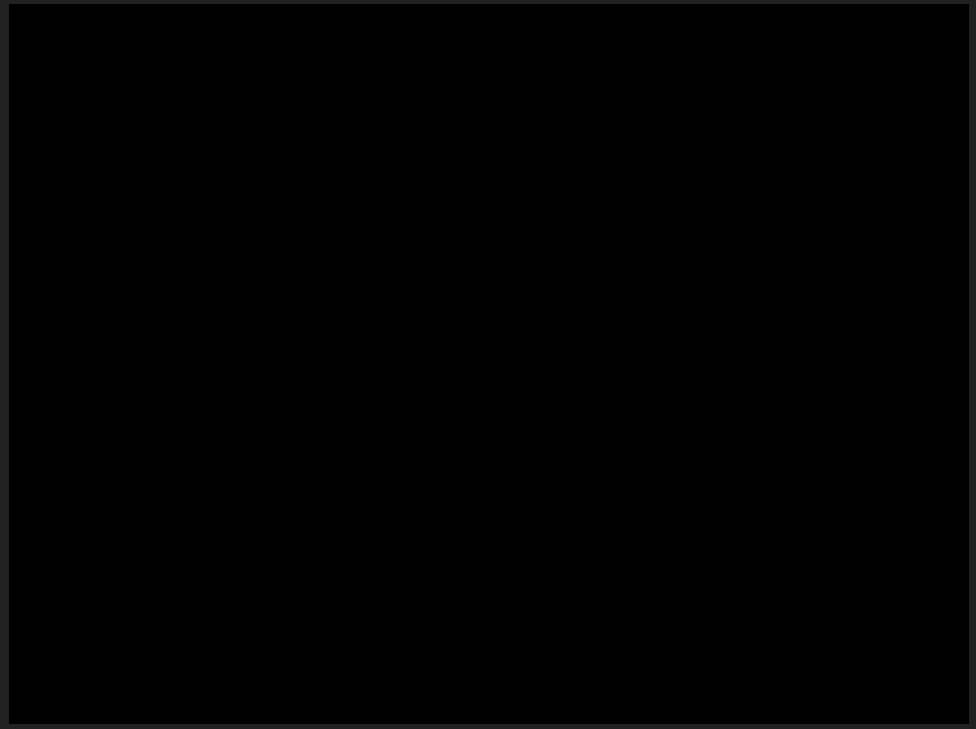

After our experiment, we answered all three of these questions.

---

1. Which lifestyle choices are most associated with a higher severity of cancer?
   Our model predicted occupational hazards, air pollution, and smoking as the three most relevant lifestyle factors.

2. How can we most accurately predict cancer severity using machine learning?
   We can try to run the algorithm again without some of the features that had a minimal impact on this models predictions.

3. Which machine learning model predicts the severity of cancer most accurately?
   We found that support vector machine (SVM) gave us the most accurate results.

Questions?